COMP90038

Algorit plexity

Lecture 14: Transfo er

(with thanks to Harald Sønde hael Kirley)

Andres Munoz-Acosta

munoz.m@unimelb.edu.au

Peter Hall Building G.83

Exercise: Finding Anagrams

- An **anagram** of a word *w* is a word which uses the same letters as *w* but in a different order.

- Example: 'ate', 'tea' an

- Example: 'post', 'spot', 'pots' and 'tops' are anagrams.

- Example: 'garner' and 'ranger' are anagrams.

Exercise: Finding Anagrams

- You are given a very long list of words:

{health, revolution, fool se, traverse, anger, ranger,

- Devise an algorithm to find all anagrams in the list.

Transform and Conquer

- **Instance simplificatio**

- **Representational cha**

- Problem reduction

Instance Simplification

- General principle: Try to make the problem easier through some type of pre-processing, typically sorting.

- We can pre-sort input ple
  - finding the **median**
  - **uniqueness checking**
  - finding the **mode**

Uniqueness Checking, Brute-Force

- The problem:
- Given an unsorted array A[0]…A[n-1], is A[i]≠A[j] whenever i≠j?

- The obvious approach is

- **What is the complexity of this?**

Uniqueness Checking, with Pre-sorting

- Sorting makes the problem easier:

- **What is the complexity of this?**

Exercise: Computing a Mode

- A **mode** is a list or array element which occurs most frequently in the list/array. For example, in

[42, 78, 1 ... 98, 42, 33]

the elements 13 and 42 are modes.

- The problem:
- Given array A, find a mode.

- Discuss a brute-force approach vs a pre-sorting approach.

Mode Finding, with Pre-sorting

- Again, after sorting, the rest takes linear time.

Searching, with Pre-sorting

- The problem:
- Given unsorted array A, find item *x* (or determine that it is absent).

- Compare these two ap
  - Perform a sequential search
  - Sort, then perform binary search

- **What are the complexities of these approaches?**

Searching, with Pre-sorting

- What if we need to search for m items?
- Let us do a back-of-the envelope calculation (consider worst cases for simplicity):

- Take n = 1024 and m = 32.

- Sequential search: $m \times n = 32{,}768$.
- Sorting + binsearch: $n \log_2 n + m \times \log_2 n = 10{,}240 + 320 = 10{,}560$.
- Average-case analysis will look somewhat better for sequential search, but pre-sorting will still win.

Exercise: Finding Anagrams

- You are given a very long list of words.

- Devise an algorithm to he list.

- An approach could be to sort each word, sort the list of words, and then find the repeats…

- **What would be the time complexity?**

Exercise: Finding Anagrams

health	aehhlt	aeerstv	1
revolution	eilnoortvu	aegnr	1
foolish	fhiloos	aegnrr	1
garner	aegnrr	aegnrr	**2** (This element is an anagram)
drive			1
praise			1
traverse	aeerstv		1
anger	aegnr		1
ranger	aegnrr	deirv	1
…	…	…	…
scoop	coops	eilnoortvu	1
fall	afll	fhiloos	1
truly	lrtuy	lrtuy	1
	Sort each word	Sort the list	Find repeats

Binary Search Trees

- A **binary search tree**, or **BST**, is a binary tree that stores elements in all internal nodes, with each sub-tree satisfying the BST property:

- Let the root be *r*; then **ft subtree is smaller** than *r* and each element in the **right** **s larger** than *r*.

- (For simplicity we will assume that all keys are different.)

Binary Search Trees

- BSTs are useful for search applications. To search for k in a BST, compare against its root r. If r=k, we are done; otherwise search in the left or right sub-tree, according as k<r or k>r.

- If a BST with n elements is "reasonably" balanced, search involves, in the worst case, $\Theta(\log n)$ comparisons.

Binary Search Trees

- If the BST is not well balanced, search performance degrades, and may be as bad as linear search:

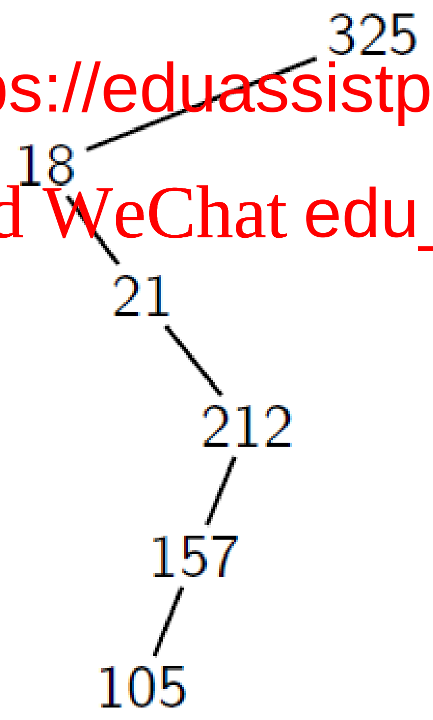

Insertion in Binary Search Trees

- To insert a new element *k* into a BST, we pretend to search for *k*.
- When the search has taken us to the fringe of the BST (we find an empty sub-tree), we insert *k* where we would expect to find it.
- Where would you insert 24?

Insertion in Binary Search Trees

- To insert a new element *k* into a BST, we pretend to search for *k*.
- When the search has taken us to the fringe of the BST (we find an empty sub-tree), we insert *k* where we would expect to find it.
- Where would you insert 24?

BST Traversal Quiz

- Performing .... l produce its elements in sorted order.

Next Up: Balancing Binary Search Trees

- To optimise the performance of BST search, it is important to keep trees (reasonably) bal

- Next we shall look at **AVL trees** and